Attention value projection: $|\gamma_{ii}|$ vs σ_i (HVP curvature along D_i)

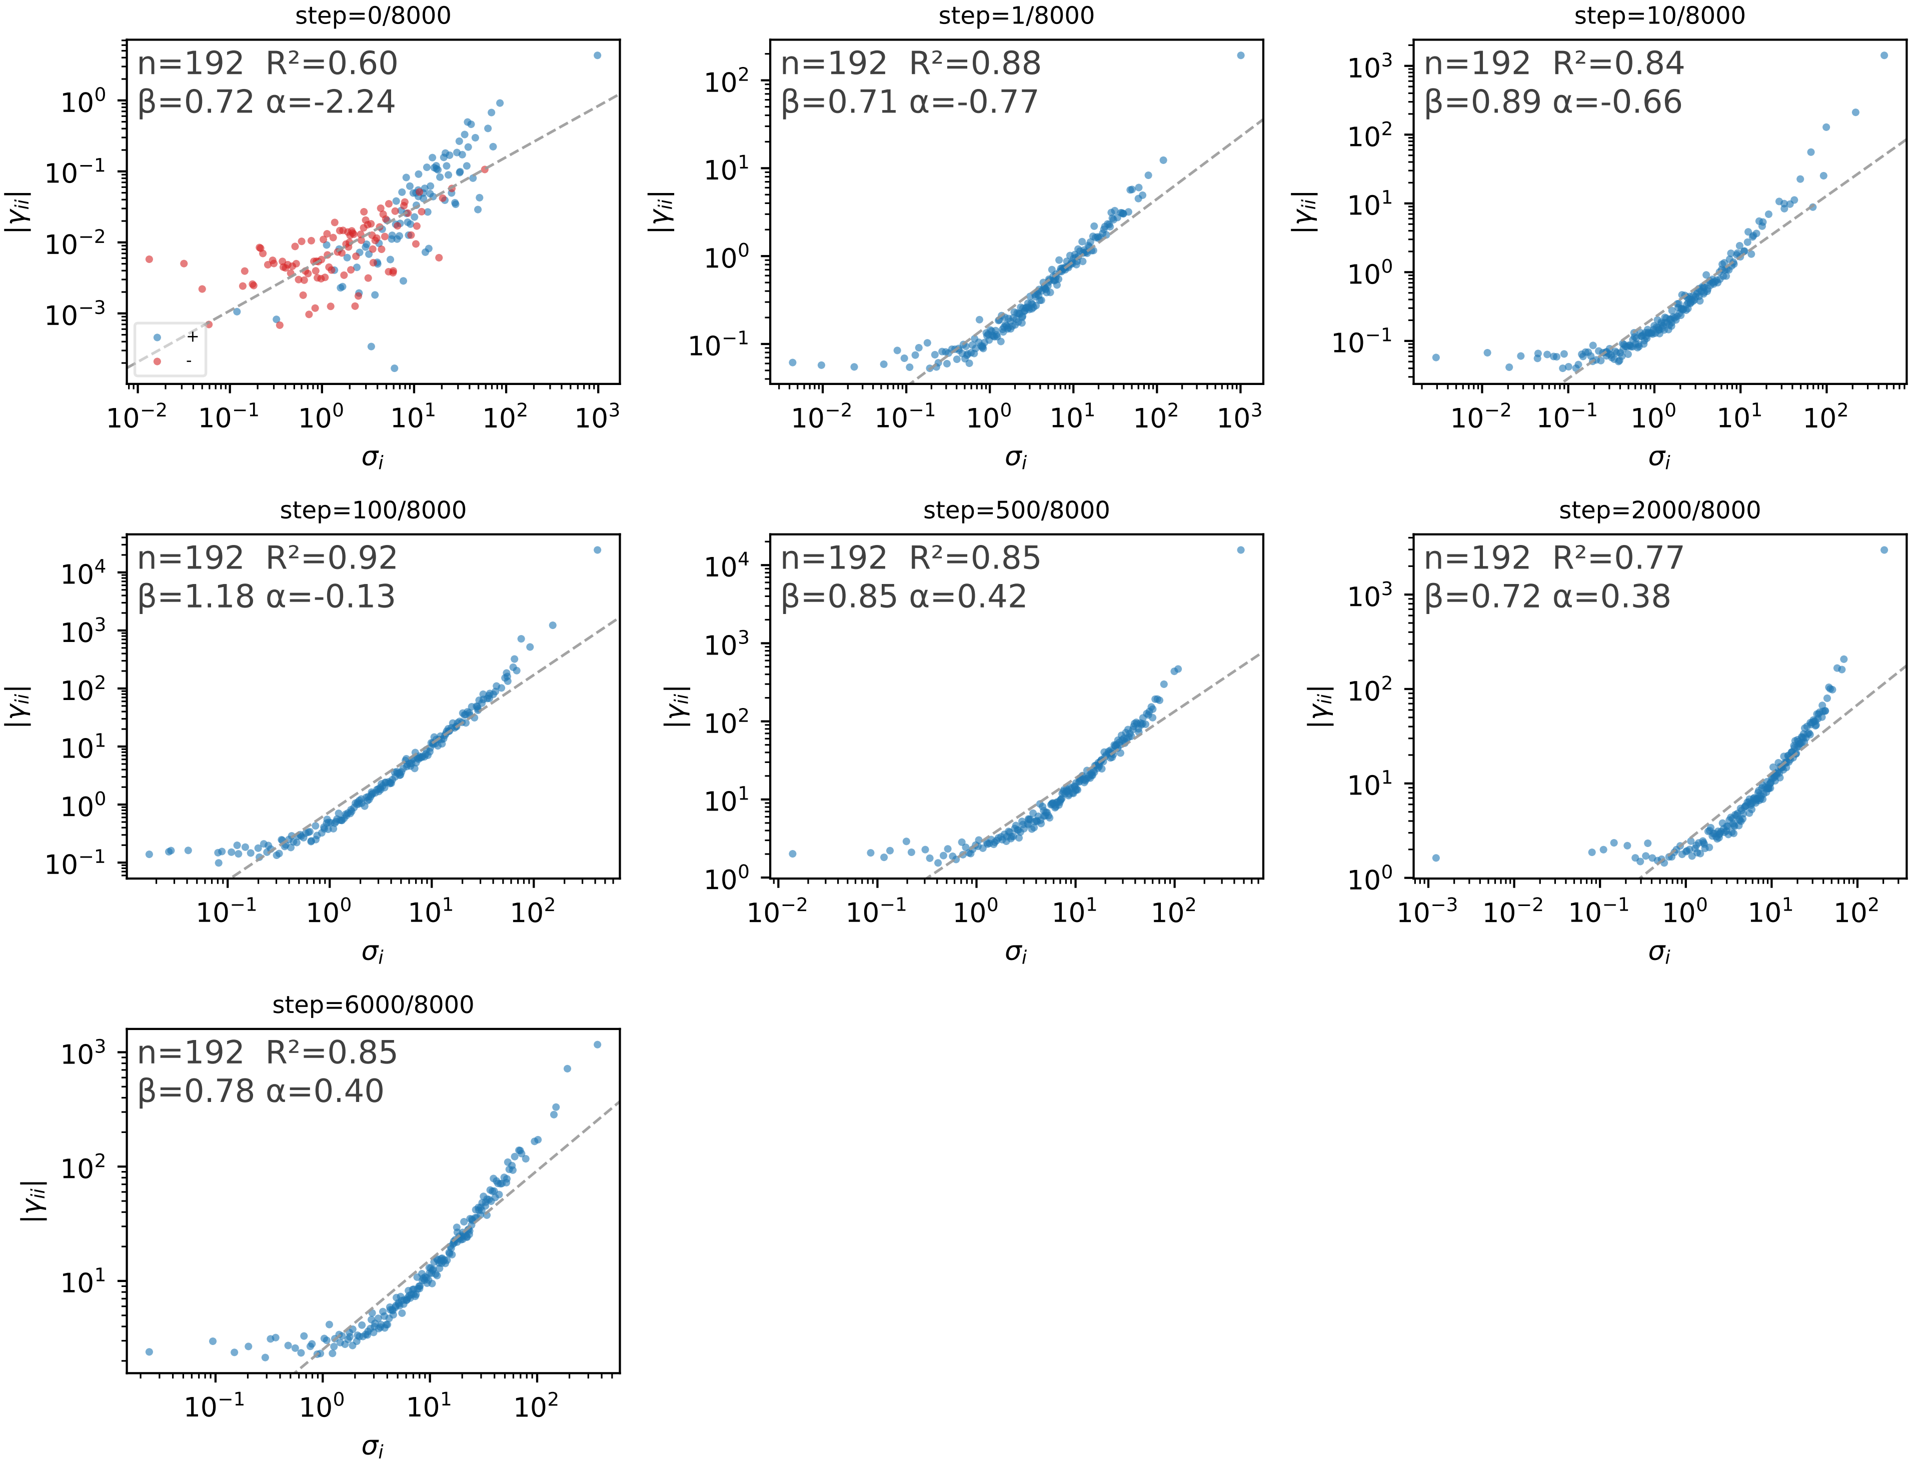

n: number of points.
R^2, β, α: least-squares fit log10(y) = β·log10(x) + α (β the fitted power-law exponent, α the intercept).

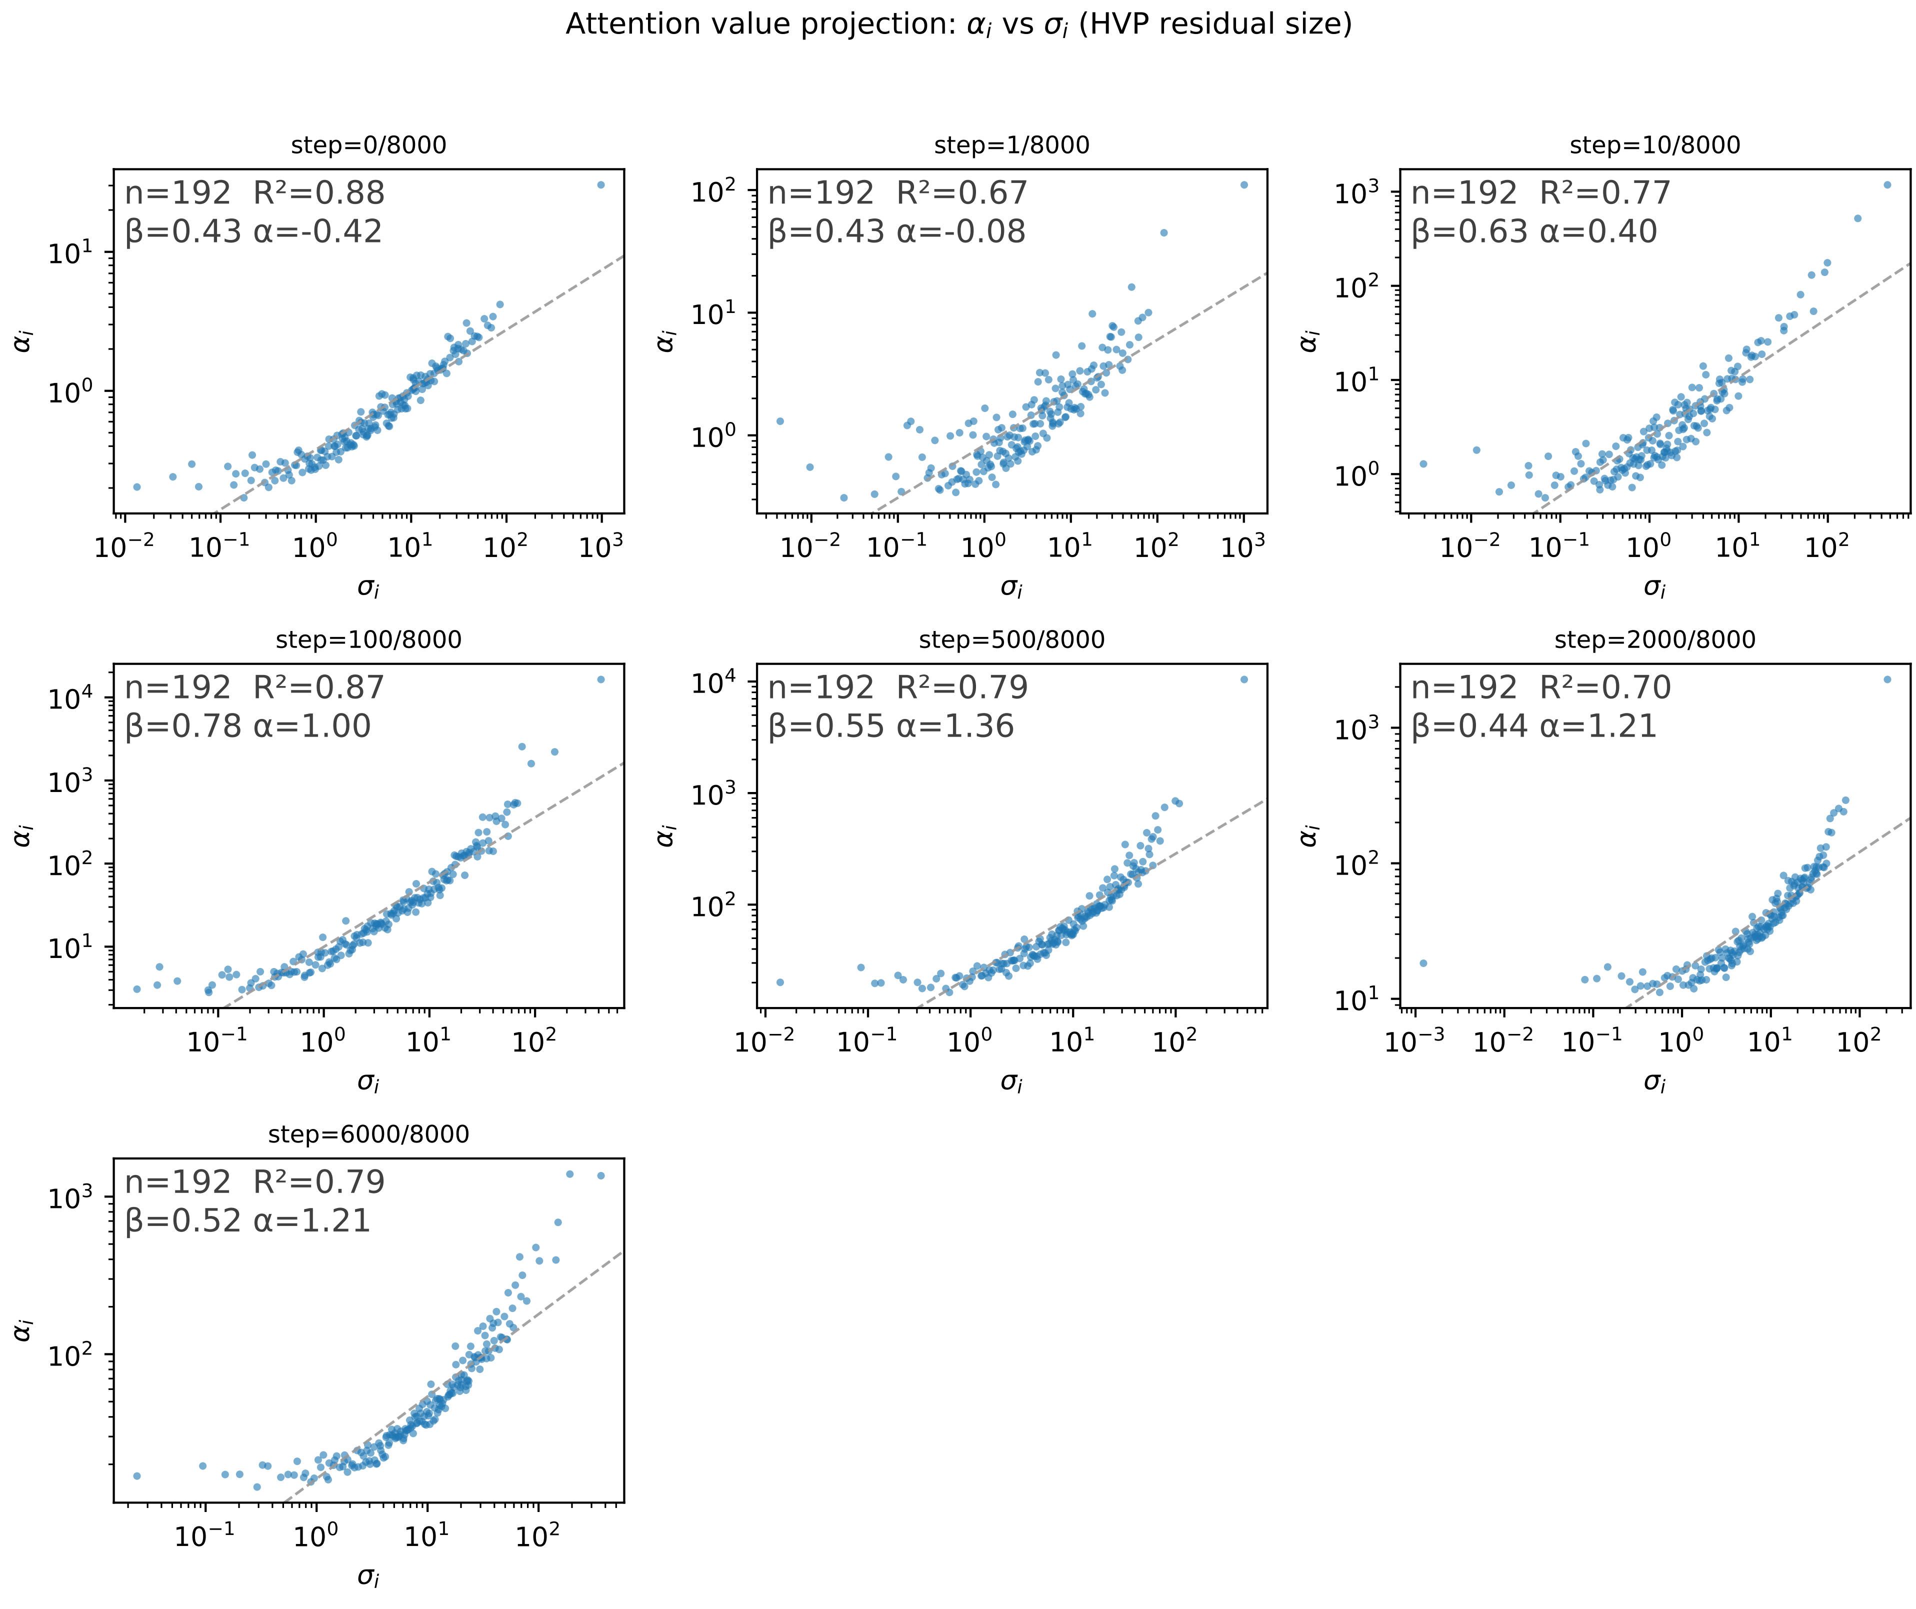

n: number of points.

R^2, β, α: least-squares fit log10(y) = β·log10(x) + α (β the fitted power-law exponent, α the intercept).

Attention value projection: $|\gamma_{ii}^{\perp}|$ vs σ_i (curvature along Q_i)

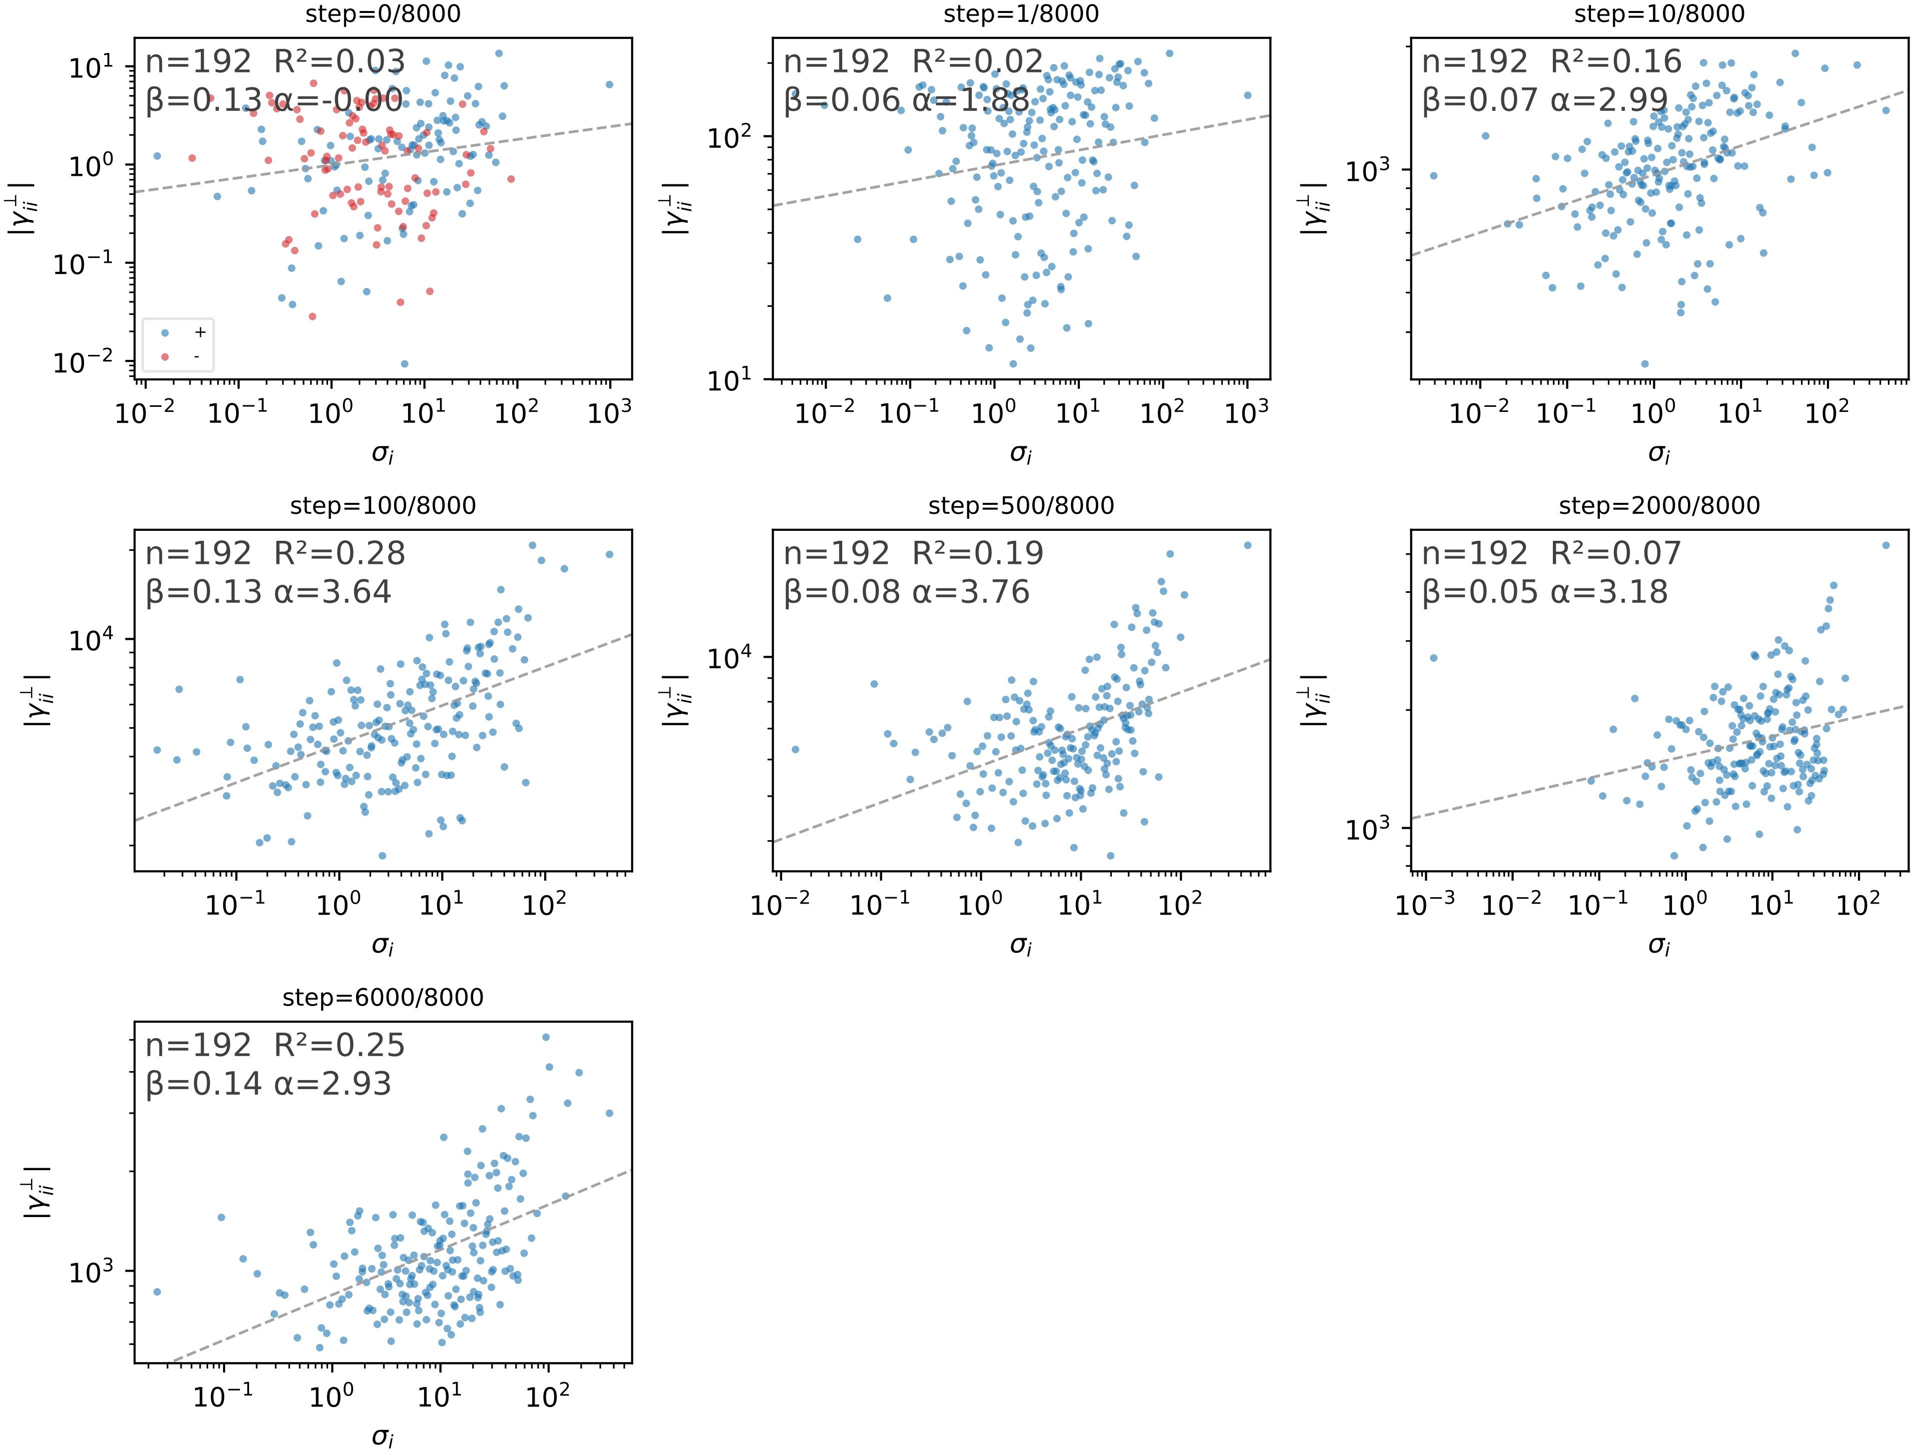

n: number of points.
R², β, α: least-squares fit log10(y) = β·log10(x) + α (β the fitted power-law exponent, α the intercept).

MLP up-projection: $|\gamma_{ii}|$ vs σ_i (HVP curvature along D_i)

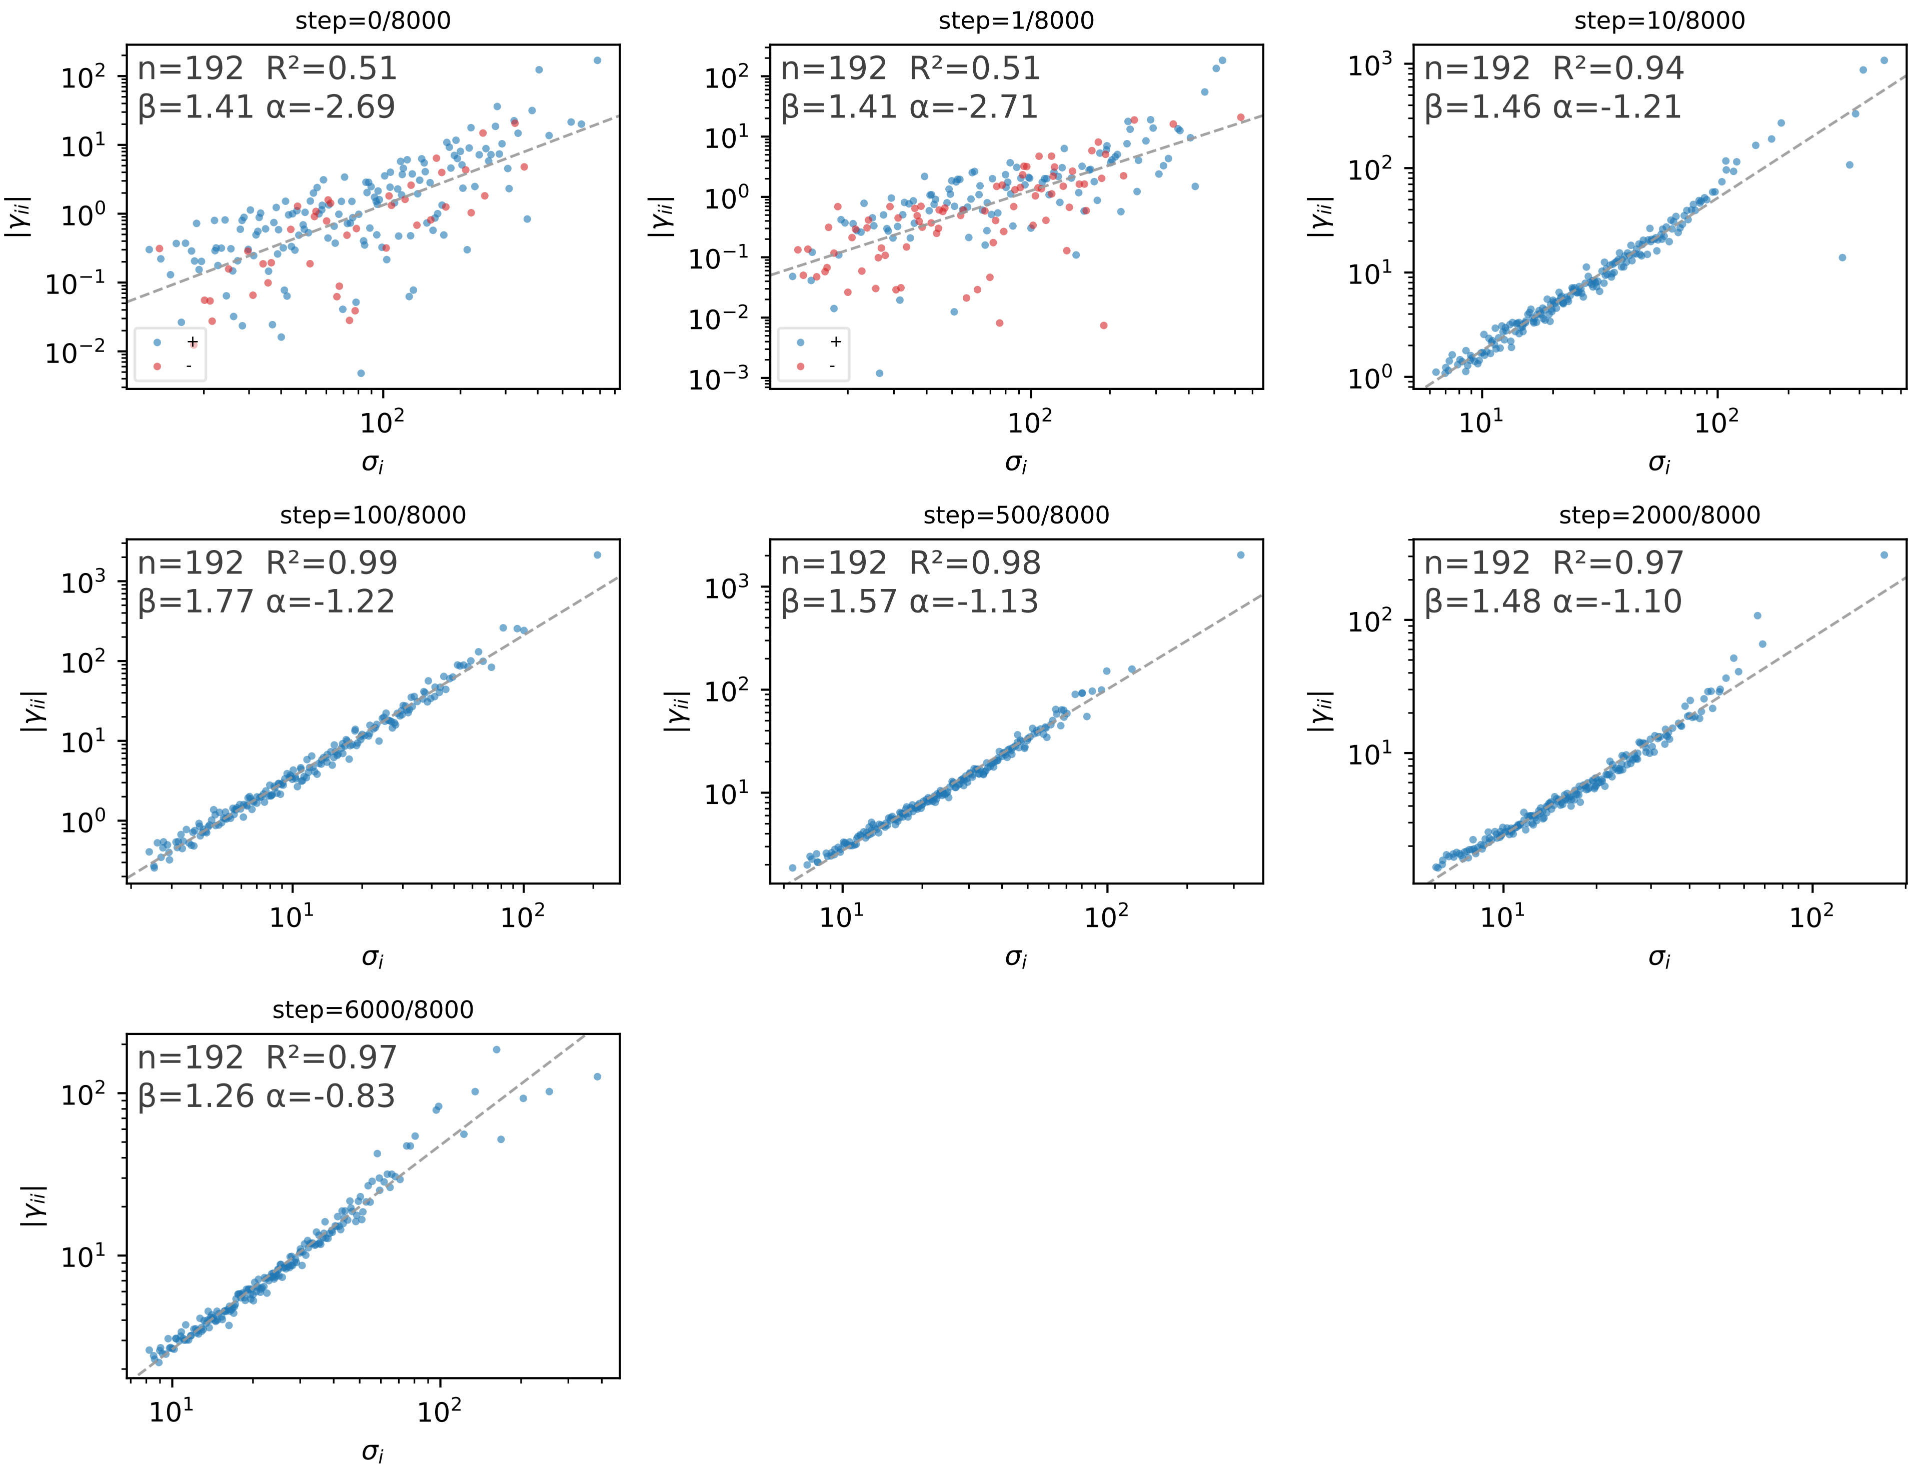

n: number of points.

R^2, β, α: least-squares fit log10(y) = β·log10(x) + α (β the fitted power-law exponent, α the intercept).

MLP up-projection: α_i vs σ_i (HVP residual size)

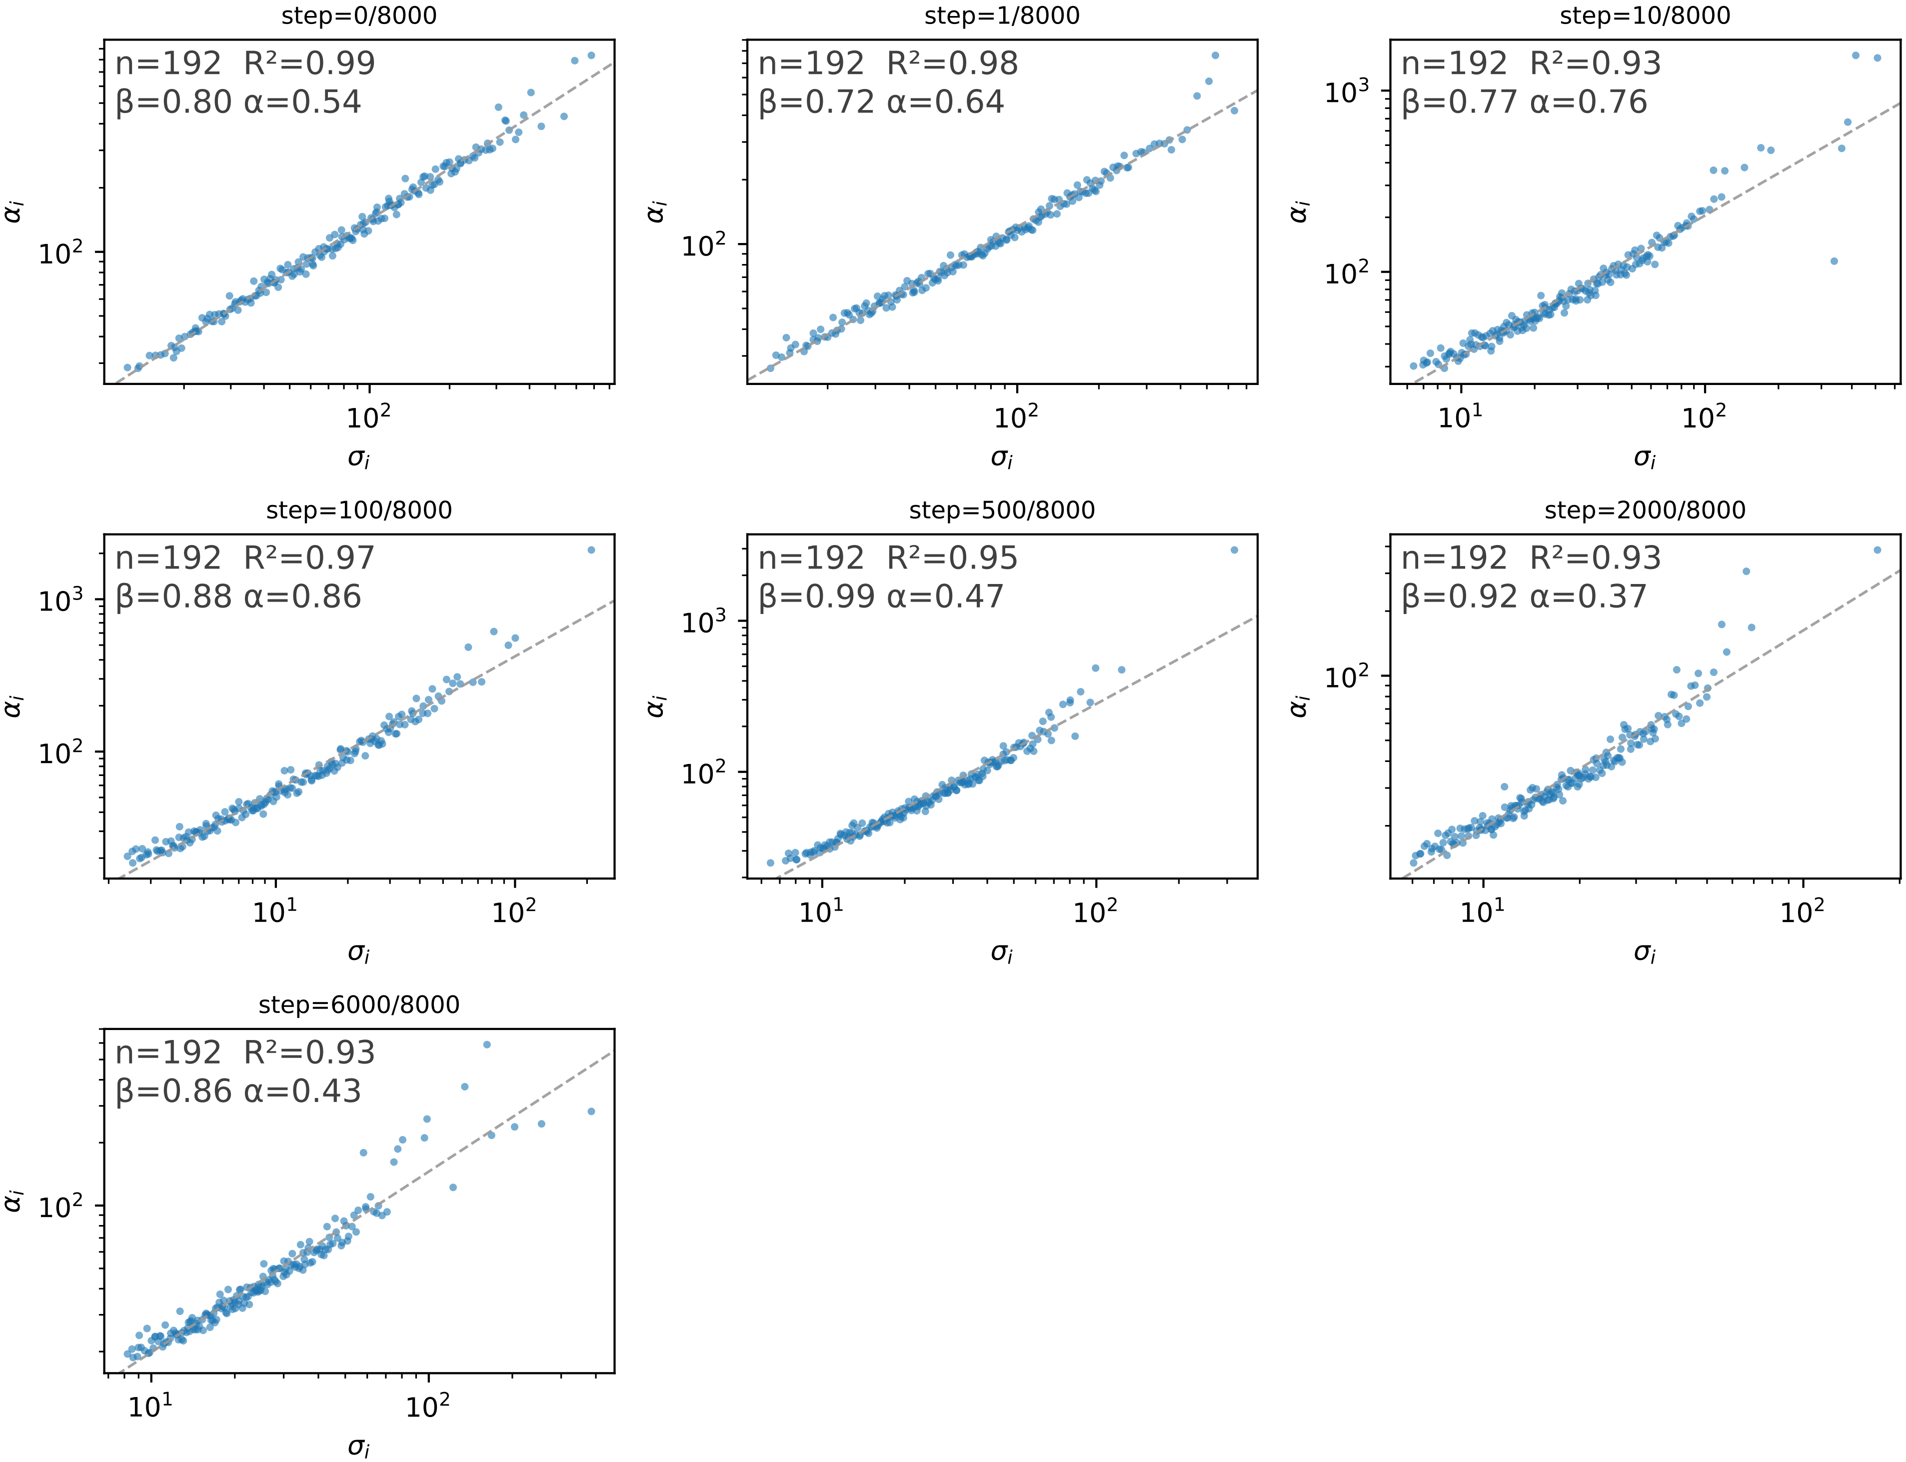

n: number of points.

R^2, β, α: least-squares fit $\log10(y) = \beta\cdot\log10(x) + \alpha$ (β the fitted power-law exponent, α the intercept).

MLP up-projection: $|\gamma_{ii}^{\perp}|$ vs σ_i (curvature along Q_i)

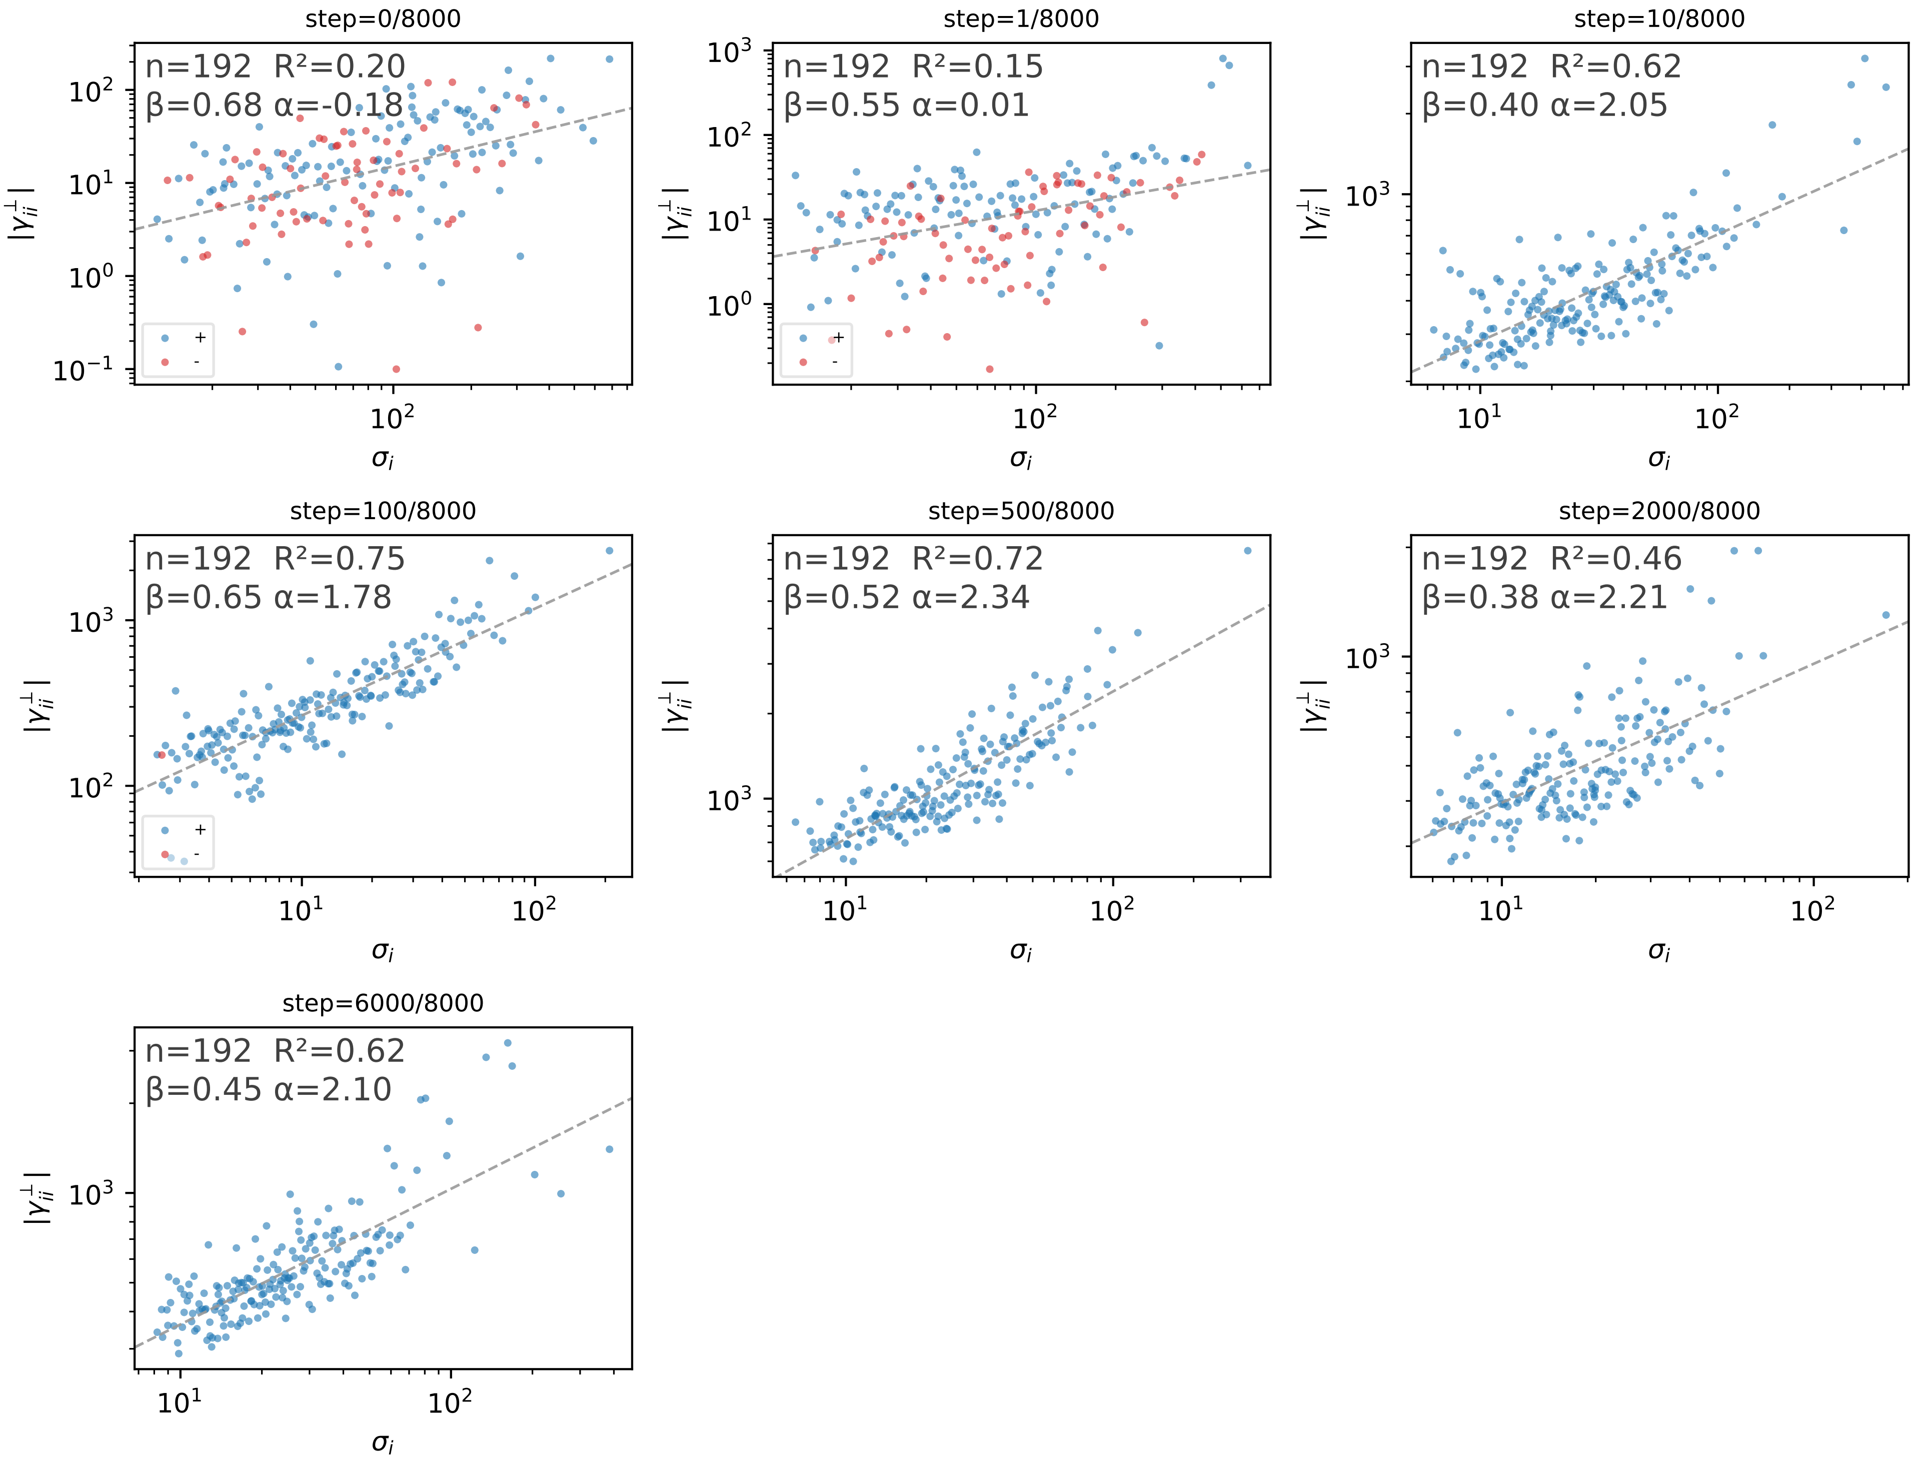

n: number of points.
R², β, α: least-squares fit log10(y) = β·log10(x) + α (β the fitted power-law exponent, α the intercept).