


Attention value projection: $|\gamma_{ii}|$ vs σ_i (HVP curvature along D_i)




n : number of points.


R^2 , β , α : least-squares fit $\log_{10}(y) = \beta \cdot \log_{10}(x) + \alpha$ (β the fitted power-law exponent, α the intercept).

Attention value projection: α_i vs σ_i (HVP residual size)


step=0/8000




step=1/8000




step=10/8000




step=100/8000




step=500/8000



step=2000/8000




step=6000/8000



n: number of points.

R^2 , β , α : least-squares fit $\log_{10}(y) = \beta \cdot \log_{10}(x) + \alpha$ (β the fitted power-law exponent, α the intercept).


Attention value projection: $|y_{ii}^\perp|$ vs σ_i (curvature along Q_i)



n: number of points.

R^2 , β , α : least-squares fit $\log_{10}(y) = \beta \cdot \log_{10}(x) + \alpha$ (β the fitted power-law exponent, α the intercept).


MLP up-projection: $|\gamma_{ii}|$ vs σ_i (HVP curvature along D_i)



n: number of points.

R^2 , β , α : least-squares fit $\log_{10}(y) = \beta \cdot \log_{10}(x) + \alpha$ (β the fitted power-law exponent, α the intercept).


MLP up-projection: α_i vs σ_i (HVP residual size)



n : number of points.

R^2 , β , α : least-squares fit $\log_{10}(y) = \beta \cdot \log_{10}(x) + \alpha$ (β the fitted power-law exponent, α the intercept).

MLP up-projection: $|y_{ii}^\perp|$ vs σ_i (curvature along Q_i)



n: number of points.

R^2 , β , α : least-squares fit $\log_{10}(y) = \beta \cdot \log_{10}(x) + \alpha$ (β the fitted power-law exponent, α the intercept).