



Attention value projection: $|y_{ij}|$ vs σ_i , colored by step since fork



n: number of points.

R^2 , β , α : least-squares fit $\log_{10}(y) = \beta \cdot \log_{10}(x) + \alpha$ (β the fitted power-law exponent, α the intercept).

MLP up-projection: $|Y_{ii}|$ vs σ_i , colored by step since fork



n: number of points.

R^2 , β , α : least-squares fit $\log_{10}(y) = \beta \cdot \log_{10}(x) + \alpha$ (β the fitted power-law exponent, α the intercept).